

Sai Nageswar Satchidanand

Senior Software Engineer / Applied Scientist



Contact Information:

📍 Hyderabad, India

✉️ sainageswar@alumni.iitm.ac.in

☎️ +91 70223 74569

🌐 [linkedin.com/in/sainageswar](https://www.linkedin.com/in/sainageswar)

🐙 github.com/SaiNageswarS

Summary

Senior Software Engineer with deep expertise in **Agents, LLMs, Search, Distributed Systems, and ML Pipelines**. Proven experience building **large-scale, low-latency systems** for production AI, including turn based agents, semantic search, caching, and observability. Strong background in **causal inference, deep learning, big data**, and **cloud-native platforms**. Passionate about system design, applied ML, and developer productivity.

Work Experience

Senior Applied Scientist — Microsoft

Core AI & Search (Copilot Search)

May 2022 – Present

- **Architected and deployed turn-based agent systems** serving **10M+ queries daily**, improving answer satisfaction score by **35%** through optimized SLM planning and LLM reasoning pipelines.
- **Built production-scale offline pipelines** generating LLM-based titles, captions, and web answers for **Bing's top 1M queries**.
- **Developed comprehensive embedding model experiments** and extended HNSW-based retrieval from existing QnA systems to captions and content components, improving **content coverage by 40%** and achieving **~30ms P95 retrieval latency** across **150M+ indexed items**.
- **Engineered low-latency online LLM pipelines** for Captions serving long-tail queries by **architecting pre-ranking and post-ranking system design**, triggering LLM generations during pre-ranking with **delayed cache reads** before rendering, maintaining **~200ms page render time** while serving **2M+ daily requests**.
- **Implemented comprehensive observability framework** across 15+ pipeline stages, enabling **real-time performance optimization** and reducing system downtime by **90%**.

- **Architected and deployed end-to-end experimentation infrastructure** from ground up, including **experiment orchestration pipelines**, **interactive Gradio-based experiment interfaces**, **secure production service integrations**, and **network-isolated access controls**, supporting **100+ concurrent A/B experiments** while maintaining **secure API authentication** across **20+ microservices** with **zero security incidents**.
- **Led prompt engineering initiatives** improving grounded answer accuracy by **25%** through advanced chain-of-thought prompting and web data integration.

Senior Software Engineer — Microsoft

Substrate Platform, Office 365

Jan 2021 – May 2022

- **Architected and implemented a Kubernetes-for-Kubernetes control plane** managing **50+ multi-region clusters** serving **100M+ Office 365 users** with **99.9% uptime**.
- **Architected reconciliation systems** maintaining desired vs actual state across clusters, reducing configuration drift by **95%** and deployment failures by **80%**.
- **Designed and secured end-to-end infrastructure** managing Key Vault, logging sidecars, monitoring, and API-to-API auth for **entire Office 365 substrate**, eliminating **100+ manual intervention points**.
- **Created reusable typesafe spec framework** adopted by **5+ teams**, replacing error-prone YAML configurations and reducing deployment bugs by **70%**.
- **Built composite inventory abstraction** enabling real-time topology queries across **multiple storage backends**, improving system visibility and reducing troubleshooting time by **60%**.
- **Led technical design and implementation** for critical platform components, mentoring **3 junior engineers** and establishing coding standards for the team.

SDE II — Amazon

Consumer Behavior Analytics

Jun 2017 – Dec 2020

- Applied **causal inference** to measure incremental impact of ads on purchases and high-value user actions (Prime signup, Video/Music usage).
- Trained **DNN models** mapping control and treatment variables to user outcomes.
- Built **feature engineering pipelines** to capture correlations between user actions.
- Modeled user-action similarity using **graph representations (Jaccard index)** and applied **spectral analysis**.
- Processed **terabytes of data** using distributed Spark pipelines for feature transformation and inference.
- Implemented **regular model retraining**, weight regularization, and experimented with **sequential models** for campaign effects.

SDE II — Codenation / Trilogy Innovations

May 2016 – Jun 2017

- Mined large-scale source-code repositories and built **AST embeddings** for duplicate and similar code detection.
- Developed **auto-refactoring tools** using Eclipse JDT and Clone Digger.
- Built an **organization-wide executive dashboard** tracking developer productivity using statistical analysis.

- Deployed services using **Docker-based orchestration**.

Software Development Engineer — Gozoomo

Jul 2015 – May 2016

- Built a **peer-to-peer used-car marketplace** with scalable microservices.
- Auto-generated **typesafe client/server code** using Protobuf.
- Implemented **secure API-to-API authentication**, monitoring, and alerting.
- Developed a **custom chat framework** using Firebase and Angular.

Software Development Engineer — Codenation / Trilogy Innovations

Jul 2014 – Jun 2015

- Worked on **enterprise analytics and data warehousing systems** using Microsoft SQL Server.
- Modernized legacy systems with refactoring, unit tests, and mocking.
- Added **CI/CD pipelines** with integration testing.

Education

M.Tech, Computer Science Engineering

Indian Institute of Technology Madras

2012 – 2014

- Research: *Semi-Supervised Multi-View Multi-Relation Collective Inference using Hypergraphs*
- Published at **IJCAI 2015**

B.Tech, Computer Science Engineering

National Institute of Science and Technology

2006 – 2010

- Built a **scientific mathematics tool** parsing free-form equations and solving ODEs with graph visualization.

Publications

Extended Discriminative Random Walk: A Hypergraph Approach to Multi-View Multi-Relational Transductive Learning

IJCAI 2015

Sai Nageswar Satchidanand, Harini Ananthapadmanaban, Balaraman Ravindran

<https://dl.acm.org/doi/abs/10.5555/2832747.2832778>

Open Source Projects

go-api-boot

<https://github.com/SaiNageswarS/go-api-boot>

Production-grade **Go gRPC microservice framework** with zero-config HTTPS, MongoDB ODM (vector search), Temporal workflows, and cloud abstractions (Azure/GCP).

GraphMind

<https://github.com/SaiNageswarS/GraphMind>

Semantic code intelligence platform that builds **RDF knowledge graphs across repositories**, enabling LLM-powered multi-repo code generation.

agent-boot

<https://github.com/SaiNageswarS/agent-boot>

Multi-tenant **RAG system** integrating Claude with domain-specific knowledge using Go + Python ML pipelines.

GeoSpatialAnalysis

<https://github.com/SaiNageswarS/GeoSpatialAnalysis>

Scalable remote-sensing pipeline using **Temporal + GDAL** for satellite imagery (MODIS/FAPAR) processing.

Skills

- **Languages:** Go, C++, Java, Python, C#, SQL
- **ML/AI:** LLMs, Embeddings, Vector Search, Causal Inference, Deep Learning
- **Systems:** Distributed Systems, gRPC, Kubernetes, Temporal, Spark
- **Cloud:** Azure, GCP, Docker, IaC
- **Databases:** MongoDB, SQL Server, NoSQL